

Robustness -1:

Theory and rethinking

Yunzhen Feng, Haocheng Ju, Zehao Wang, Haotong Yang, Pu Yang

School of Mathematical Science, Peking University

The Landscape of Robust Models

The Landscape of Robust Models

**Adversarial Robustness through Local
Linearization**

Table of contents

1. The Landscape of Robust Models

Adversarial Robustness through Local Linearization

2. Certified Robustness

3. The Generalization of Robust Model

4. Rethinking Different Robustness

Gradient Obfuscation and Robustness ¹

- Landscape for failed defense and clean training:
- Highly non-linear in the vicinity.
- Adversarial training on stronger attack make surface more linear.
- How to train with weak adversary and prevent gradient obfuscation?




Figure 1: The landscape surface

¹Qin, Chongli, et al. "Adversarial robustness through local linearization." Advances in Neural Information Processing Systems. 2019.

Local Linearity and Robustness

- Local Linearity Measure:

$$\gamma(\epsilon, x) = \max_{\delta \in B(\epsilon)} |\ell(x + \delta) - \ell(x) - \delta^T \nabla_x \ell(x)|.$$





Figure 2: The change of γ w.r.t. different training settings.

Induced ‘Local Linearity Regularizer’

- By definition of $\gamma(\epsilon, \mathbf{x})$:

$$\ell(\mathbf{x} + \delta) \leq \ell(\mathbf{x}) + |\delta^T \nabla_{\mathbf{x}} \ell(\mathbf{x})| + \gamma(\epsilon, \mathbf{x}).$$

- For small δ , the middle part is small
- Specifically, it can be controlled by $\gamma(\epsilon, \mathbf{x})$ for quadratic loss or cross-entropy.
- Objective function:

$$L(\mathcal{D}) = \mathbb{E}_{\mathcal{D}}[\ell(\mathbf{x}) + \underbrace{\lambda\gamma(\epsilon, \mathbf{x}) + \mu |\delta_{LLR}^T \nabla_{\mathbf{x}} \ell(\mathbf{x})|}_{LLR}],$$

where $\delta_{LLR} = \operatorname{argmax}_{\delta \in B(\epsilon)} |\ell(\mathbf{x} + \delta) - \ell(\mathbf{x}) - \delta^T \nabla_{\mathbf{x}} \ell(\mathbf{x})|$ and $\gamma(\epsilon, \mathbf{x}) = |\ell(\mathbf{x} + \delta_{LLR}) - \ell(\mathbf{x}) - \delta_{LLR}^T \nabla_{\mathbf{x}} \ell(\mathbf{x})|$.

Algorithm 1 Local Linearization of Network

Require: Training data $X = \{(x_1, y_1), \dots, (x_N, y_N)\}$. Learning rate lr and batch size for training b and number of iterations N . Number of iterations for inner optimization M and step size s and network architecture parameterized by θ .

- 1: Initialize variables θ .
 - 2: **for all** $i \in \{0, 1, \dots, N\}$ **do**
 - 3: Get mini-batch $B = \{(x_{i_1}, y_{i_1}), \dots, (x_{i_b}, y_{i_b})\}$.
 - 4: Calculate loss wrt to minibatch $L_B = \frac{1}{b} \sum_{j=1}^b \ell(x_{i_j}, y_{i_j})$.
 - 5: Initialize initial perturbation δ uniformly in the interval $[-\epsilon, \epsilon]$.
 - 6: **for all** $j \in \{0, 1, \dots, M\}$ **do**
 - 7: Calculate $g = \frac{1}{b} \sum_{t=1}^b \nabla_{\delta} g(\delta; x_{i_t}, y_{i_t})$ at δ .
 - 8: Update $\delta \leftarrow \text{Proj}(\delta - s \times \text{Optimizer}(g))$
 - 9: **end for**
 - 10: Compute objective $L = L_B + 1/b \sum_{j=1}^b (\lambda g(\delta; x_{i_j}, y_{i_j}) + \mu \|\delta^T \nabla_x l(x)\|)$
 - 11: $\theta \leftarrow \theta - lr \times \text{Optimizer}(\nabla_{\theta} L)$
 - 12: **end for**
-

Note $g(\delta; x, y) = \ell(x + \delta; y) - \ell(x; y) - \delta^T \nabla_x \ell(x; y)$.

Results

CIFAR-10: Wide-ResNet-28-8 (8/255)				
Methods	Nominal	FGSM-20	Untargeted	Multi-Targeted
Attack Strength		Weak	Strong	Very Strong
ADV[16]	87.25%	48.89%	45.92%	44.54%
CURE[19]	80.76%	39.76%	38.87%	37.57%
ADV(S)	85.11%	56.76%	53.96%	48.79%
CURE(S)	84.31%	48.56%	47.28%	45.43%
TRADES(S)	87.40%	51.63	50.46%	49.48%
LLR (S)	86.83%	54.24%	52.99%	51.13%
CIFAR-10: Wide-ResNet-40-8 (8/255)				
ADV(R)	85.58%	56.32%	52.34%	46.89%
TRADES(R)	86.25%	53.38%	51.76%	50.84%
ADV(S)	85.27%	57.94%	55.26%	49.79%
CURE(S)	84.45%	49.41%	47.69%	45.51%
TRADES(S)	88.11%	53.03%	51.65%	50.53%
LLR (S)	86.28%	56.44%	54.95%	52.81%

Figure 3: Results on Cifar-10


Results

ImageNet: ResNet-152 (4/255)				
Methods	PGD steps	Nominal	Untargeted	Random-Targeted
Accuracy				Success Rate
ADV	30	69.20%	39.70%	0.50%
DENOISE	30	69.70%	38.90%	0.40%
LLR	2	72.70%	47.00%	0.40%
ImageNet: ResNet-152 (16/255)				
ADV [28]	30	64.10%	6.30%	40.00%
DENOISE [28]	30	66.80%	7.50%	38.00%
LLR	10	51.20%	6.10%	43.80%


Figure 4: Results on ImageNet

- Trained on 128 TPUv3 cores for radius 4/255.
- ADV networks takes 36 hours for 110 epochs
- LLR networks takes 7 hours for 110 epochs

Results




(a) CIFAR-10 (8/255)




(b) ImageNet (4/255)


Figure 5: Robustness Drop is smaller




(a) ADV-1



(b) LLR-1



(c) ADV-2



(d) LLR-2

Certified Robustness

Certified Robustness

**Adversarial Training and Provable
Defenses: Bridging the Gap**

Table of contents

1. The Landscape of Robust Models

2. Certified Robustness

Adversarial Training and Provable Defenses: Bridging the Gap

3. The Generalization of Robust Model

4. Rethinking Different Robustness

The most secure defense: certified robustness ⁵

- Attacks vs Defense: ongoing war
- Can we certify the vicinity of the data?
- Linear Relaxation ²
- Interval Bound Propagation ³
- Randomized Smoothing (certified with probability) ⁴

²Zhang, Huan, et al. "Efficient Neural Network Robustness Certification with General Activation Functions." Advances in Neural Information Processing Systems (2018): 4939-4948.

³Gowal, Sven, et al. "On the Effectiveness of Interval Bound Propagation for Training Verifiably Robust Models.." arXiv: Learning (2018).

⁴Cohen, Jeremy, Elan Rosenfeld, and J. Zico Kolter. "Certified Adversarial Robustness via Randomized Smoothing." International Conference on Machine Learning (2019): 1310-1320.

⁵Balunovic, Mislav, and Martin Vechev. "Adversarial training and provable defenses: Bridging the gap." International Conference on Learning Representations. 2020.

- Neural networks in layerwise:

$$h_{\theta} = h_{\theta}^k \circ h_{\theta}^{k-1} \dots \circ h_{\theta}^1 \text{ and } h_{\theta}^i : \mathbb{R}^{d_{i-1}} \rightarrow \mathbb{R}^{d_i}.$$

- Robustness (even for other properties):

$$c^T h_{\theta}(\mathbf{x}') + d < 0, \forall \mathbf{x}' \in \mathbb{S}_0(\mathbf{x}),$$

where $\mathbb{S}_0(\mathbf{x}) = \{\mathbf{x}' \in \mathbb{R}^{d_0}, \|\mathbf{x} - \mathbf{x}'\|_{\infty} < \epsilon\}$, c, d are specific vectors.

- Perturbation set: $\mathbb{S}_i(\mathbf{x}) = h_{\theta}^i(\mathbb{S}_{i-1}(\mathbf{x}))$.
- Certification via convex relaxation:
- Find convex set $\mathbb{C}_i(\mathbf{x})$: $\mathbb{S}_i(\mathbf{x}) \subseteq \mathbb{C}_i(\mathbf{x}), \forall i \in [k]$.
- Directly prove the robustness property for $\mathbb{C}_k(\mathbf{x})$.

Specifically

- $\mathbb{S}_0(\mathbf{x}) = \mathbb{C}_0(\mathbf{x})$
- All convex set can be represented as

$$\mathbb{C}_l(\mathbf{x}) = \{\mathbf{a}_l + \mathbf{A}_l \mathbf{e} | \mathbf{e} \in [-1, 1]^{m_l}\},$$

where \mathbf{a}_l is the center.

- For $\mathbb{C}_0(\mathbf{x})$, $\mathbf{a}_0 = \mathbf{x}$ and $\mathbf{A}_0 = \epsilon \mathbf{I}_{d_0}$.
- How to propagate small enough convex set $\mathbb{C}_l(\mathbf{x})$?
- For a point in convex set $\mathbf{x}'_i = \mathbf{a}_i + \mathbf{A}_i \mathbf{e}$
- Convolutional and fully connected layer: Already Linear:

$$h_{\theta}^{i+1}(\mathbf{x}'_i) = \mathbf{W}_{i+1} \mathbf{a}_i + \mathbf{b}_{i+1} + \mathbf{W}_{i+1} \mathbf{A}_i \mathbf{e}.$$

- For ReLU activation: upper bound and lower bound

$$l_{i,j} = a_{i,j} - \sum_{k=1}^{m_i} |A_{i,j,k}|, \quad u_{i,j} = a_{i,j} + \sum_{k=1}^{m_i} |A_{i,j,k}|.$$

How to Train the Model?


- Like randomized smoothing, certification only provides a way to verify.
- How to train a model that has good certified robustness?
- Certified Robustness vs Adversarial Training

$$\min_{\theta^{l+1:k}} \mathbb{E}_{(\mathbf{x}, y) \sim D} \max_{\mathbf{x}'_l \in \mathbb{C}_l(\mathbf{x})} \mathcal{L}(h_{\theta}^{l+1:k}(\mathbf{x}'_l), y, \theta),$$

Certified Robustness find upper bound for the inner maximum while adversarial training find large enough lower bound.

- When the certification tends to fail?
- There exists $\mathbf{x}'_1 \in \mathbb{C}_1(\mathbf{x}) \setminus \mathbb{S}_1(\mathbf{x})$ that is an adversarial example.
- Add these adversarial examples into the training as adv train!

Hierarchical Adversarial Training



Algorithm 1: Convex layerwise adversarial training via convex relaxations

Data: k -layer network h_{θ} , training set $(\mathcal{X}, \mathcal{Y})$, learning rate η , step size α , inner steps n

Result: Certifiably robust neural network h_{θ}

```

1 for  $l \leq k$  do
2   for  $i \leq n_{\text{epochs}}$  do
3     Sample mini-batch  $\{(x_1, y_1), (x_2, y_2), \dots, (x_b, y_b)\} \sim (\mathcal{X}, \mathcal{Y})$ ;
4     Compute convex relaxations  $C_l(x_1), C_l(x_2), \dots, C_l(x_b)$ ;
5     Initialize  $x'_1 \sim C_l(x_1), x'_2 \sim C_l(x_2), \dots, x'_b \sim C_l(x_b)$ ;
6     for  $j \leq b$  do
7       Update in parallel  $n$  times:  $x'_j \leftarrow \Pi_{C_l(x_j)}(x'_j + \alpha \nabla_{x'_j} \mathcal{L}(h_{\theta}^{l+1:k}(x'_j), y_j))$ ;
8     end
9     Update parameters  $\theta \leftarrow \theta - \eta \cdot \frac{1}{b} \sum_{j=1}^b \nabla_{\theta} \mathcal{L}(h_{\theta}^{l+1:k}(x'_j), y_j)$ ;
10  end
11  Freeze parameters  $\theta_{l+1}$  of layer function  $h_{\theta}^{l+1}$ ;
12 end
    
```


Table 1: Evaluation on CIFAR-10 dataset with L_∞ perturbation 2/255

Method	Accuracy(%)	Certified Robustness(%)
Our work	78.4	60.5
Zhang et al. (2020)	71.5	54.0
Wong et al. (2018)	68.3	53.9
Gowal et al. (2018)	70.2	50.0
Xiao et al. (2019)	61.1	45.9
Mirman et al. (2019)	62.3	45.5


Table 2: Evaluation on CIFAR-10 dataset with L_∞ perturbation 8/255

Method	Accuracy(%)	Certified Robustness(%)
Our work	51.7	27.5
Zhang et al. (2020)	54.5	30.5
Mirman et al. (2019)	46.2	27.2
Wong et al. (2018)	28.7	21.8
Xiao et al. (2019)	40.5	20.3

Results



(a) CIFAR-10, 2/255



(b) CIFAR-10, 8/255

The Generalization of Robust Model

The Generalization of Robust Model

Rademacher Complexity for Adversarially Robust Generalization^a

^aYin, Dong, Ramchandran Kannan, and Peter Bartlett. "Rademacher Complexity for Adversarially Robust Generalization." International Conference on Machine Learning. 2019.

Table of contents

1. The Landscape of Robust Models

2. Certified Robustness

3. The Generalization of Robust Model

Rademacher Complexity for Adversarially Robust Generalization⁶

4. Rethinking Different Robustness

Adversarially robust generalization

- Two different generalization:

$$R(f) := \mathbb{E}_{(\mathbf{x}, y) \in \mathcal{D}} [\ell(f(\mathbf{x}), y)], \quad \tilde{R}(f) := \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\max_{\mathbf{x}' \in \mathbb{B}_{\mathbf{x}}^a(\epsilon)} \ell(f(\mathbf{x}'), y) \right].$$




Figure 7: Two generalization gap for standard training and adversarial training

Generalization theory in two situations

- Empirical Risk

$$R_n(f) := \frac{1}{n} \sum_{i=1}^n \ell(f(\mathbf{x}_i), y_i), \quad \tilde{R}_n(f) := \frac{1}{n} \sum_{i=1}^n \max_{\mathbf{x}'_i \in \mathbb{B}_{\mathbf{x}_i}^\infty(\epsilon)} \ell(f(\mathbf{x}'_i), y_i)$$

- Population Risk

$$R(f) := \mathbb{E}_{(\mathbf{x}, y) \in \mathcal{D}} [\ell(f(\mathbf{x}), y)], \quad \tilde{R}(f) := \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\max_{\mathbf{x}' \in \mathbb{B}_{\mathbf{x}}^{\mathbf{a}}(\epsilon)} \ell(f(\mathbf{x}'), y) \right].$$

- Rademacher Complexity

$$\mathfrak{R}_{\mathcal{D}}(\ell_{\mathcal{F}}) := \frac{1}{n} \mathbb{E}_{\sigma} \left[\sup_{h \in \mathcal{F}} \sum_{i=1}^n \sigma_i h(\mathbf{z}_i) \right], \quad \mathfrak{R}_{\mathcal{D}}(\tilde{\ell}_{\mathcal{F}})$$

- Generalization Bound

$$R(f) \leq R_n(f) + 2B\mathfrak{R}_{\mathcal{D}}(\ell_{\mathcal{F}}) + 3B\sqrt{\frac{\log \frac{2}{\delta}}{2n}}$$

$$\tilde{R}(f) \leq \tilde{R}_n(f) + 2B\mathfrak{R}_{\mathcal{D}}(\tilde{\ell}_{\mathcal{F}}) + 3B\sqrt{\frac{\log \frac{2}{\delta}}{2n}}$$

Linear Classifier for Binary Classification

- $\mathcal{Y} = \{-1, +1\}$, $\mathcal{F} = \{f_{\mathbf{w}}(\mathbf{x}) : \|\mathbf{w}\|_p \leq W\}$
- Adversarial Function Class

$$\tilde{\mathcal{F}} = \left\{ \min_{\mathbf{x}' \in \mathbb{B}_{\mathbf{x}_0}(\epsilon)} y \langle \mathbf{w}, \mathbf{x}' \rangle : \|\mathbf{w}\|_p \leq W \right\}.$$

- Theorem: Suppose $\frac{1}{p} + \frac{1}{q} = 1$, there exists a universal constant $c \in (0, 1)$ such that

$$\max \left\{ \mathfrak{R}_{\mathcal{S}}(\mathcal{F}), c\epsilon W \frac{d^{\frac{1}{q}}}{\sqrt{n}} \right\} \leq \mathfrak{R}_{\mathcal{S}}(\tilde{\mathcal{F}}) \leq \mathfrak{R}_{\mathcal{S}}(\mathcal{F}) + \epsilon W \frac{d^{\frac{1}{q}}}{\sqrt{n}}.$$

- Specifically, unavoidable dimension dependence:

$$\frac{c}{2} \left(\mathfrak{R}_{\mathcal{S}}(\mathcal{F}) + \epsilon W \frac{d^{\frac{1}{q}}}{\sqrt{n}} \right) \leq \mathfrak{R}_{\mathcal{S}}(\tilde{\mathcal{F}}) \leq \mathfrak{R}_{\mathcal{S}}(\mathcal{F}) + \epsilon W \frac{d^{\frac{1}{q}}}{\sqrt{n}}.$$

Multi-class Task

- Ramp loss: $\ell(f(\mathbf{x}), y) = \phi_\gamma(f(\mathbf{x})_y - \max_{y' \neq y} f(\mathbf{x})_{y'})$

$$\phi_\gamma(t) = \begin{cases} 1 & t \leq 0 \\ 1 - \frac{t}{\gamma} & 0 < t < \gamma \\ 0 & t \geq \gamma \end{cases}.$$

- If satisfy

$$\mathbf{1} \left(y \neq \arg \max_{y' \in [K]} [f(\mathbf{x})]_{y'} \right) \leq \ell(f(\mathbf{x}), y) \leq \mathbf{1} \left([f(\mathbf{x})]_y \leq \gamma + \max_{y' \neq y} [f(\mathbf{x})]_{y'} \right)$$

- Clean generalization with margin γ :

$$\begin{aligned} & \mathbb{P}_{(\mathbf{x}, y) \sim \mathcal{D}} \left\{ y \neq \arg \max_{y' \in [K]} [f(\mathbf{x})]_{y'} \right\} \\ & \leq \frac{1}{n} \sum_{i=1}^n \mathbf{1} \left([f(\mathbf{x}_i)]_{y_i} \leq \gamma + \max_{y' \neq y_i} [f(\mathbf{x}_i)]_{y'} \right) \\ & \quad + 2\mathfrak{R}_{\mathcal{S}}(\ell_{\mathcal{F}}) + 3\sqrt{\frac{\log \frac{2}{\delta}}{2n}}. \end{aligned}$$

Linear Classifier for Multi-class

- Linear classifier : $\mathbf{W} \in \mathbb{R}^{K \times d}$, i.e., $f(\mathbf{x}) \equiv f_{\mathbf{W}}(\mathbf{x}) = \mathbf{W}\mathbf{x}$,
 $\mathcal{F} = \left\{ f_{\mathbf{W}}(\mathbf{x}) : \|\mathbf{W}^{\top}\|_{p,\infty} \leq W \right\}$.
- Standard Rademacher:

$$\mathfrak{R}_{\mathcal{S}}(\ell_{\mathcal{F}}) = \frac{2KW}{\gamma n} \mathbb{E}_{\sigma} \left[\left\| \sum_{i=1}^n \sigma_i \mathbf{x}_i \right\|_q \right].$$

- Robust Rademacher:

$$\mathfrak{R}_{\mathcal{S}}(\tilde{\ell}_{\mathcal{F}}) \leq \frac{WK}{\gamma} \left[\frac{\epsilon \sqrt{K} d^{\frac{1}{q}}}{\sqrt{n}} + \frac{1}{n} \sum_{y=1}^K \mathbb{E}_{\sigma} \left[\left\| \sum_{i=1}^n \sigma_i \mathbf{x}_i \mathbf{1}(y_i = y) \right\|_q \right] \right]$$

Binary Deep Neural Networks

- Deep Neural Networks: $f_{\mathbf{W}}(\mathbf{x}) = \mathbf{W}_L \rho(\mathbf{W}_{L-1} \rho(\cdots \rho(\mathbf{W}_1 \mathbf{x}) \cdots))$, where $\rho(x)$ is the ReLU activation function.
- Standard Rademacher for function class $\mathcal{F} = \{f_{\mathbf{W}}(\mathbf{x}) : \mathbf{W} = (\mathbf{W}_1, \mathbf{W}_2, \dots, \mathbf{W}_L), \|\mathbf{W}_h\|_{\sigma} \leq s_h, \|\mathbf{W}_h^{\top}\|_{2,1} \leq b_h, h \in [L]\} \subseteq \mathbb{R}^{\mathcal{X}}$,

$$\mathfrak{P}_S(\mathcal{F}) \leq \frac{4}{n^{3/2}} + \frac{26 \log(n) \log(2d_{\max})}{n} \|\mathbf{X}\|_F \left(\prod_{h=1}^L s_h \right) \left(\sum_{j=1}^L \left(\frac{b_j}{s_j} \right)^{2/3} \right)^{3/2}.$$

- Robust Rademacher $\tilde{\mathcal{F}} = \{(\mathbf{x}, y) \mapsto \min_{\mathbf{x}' \in \mathbb{B}_{\infty(\epsilon)}} y f_{\mathbf{W}}(\mathbf{x}') : \mathbf{W} = (\mathbf{W}_1, \mathbf{W}_2, \dots, \mathbf{W}_L), \prod_{h=1}^L \|\mathbf{W}_h\|_{\sigma} \leq r\} \subseteq \mathbb{R}^{\mathcal{X} \times \{-1, +1\}}$,

$$\mathfrak{P}_S(\tilde{\mathcal{F}}) \geq \text{cr} \left(\frac{1}{n} \|\mathbf{X}\|_F + \epsilon \sqrt{\frac{d}{n}} \right).$$

- For two layer neural networks, the problem could be solved using a surrogate loss for function class with constraints on ℓ_1 norm.

Experiments

$$\min_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \max_{\mathbf{x}'_i \in \mathbb{B}_{\mathbf{x}_i}^a(\epsilon)} \ell(f_{\mathbf{W}}(\mathbf{x}'_i), y_i) + \lambda \|\mathbf{W}\|_1.$$




Figure 8: Linear Classifier for MNIST.

Experiments




Figure 9: Linear Classifier for MNIST: the impact of input dimension on generalization gap

Experiments




Figure 10: Four layer ReLU networks for MNIST: two convolutional and two fully connected layers.

Rethinking Different Robustness

Rethinking Different Robustness

**Exploring the Landscape of Spatial
Robustness**

Table of contents

1. The Landscape of Robust Models
2. Certified Robustness
3. The Generalization of Robust Model


4. Rethinking Different Robustness

Exploring the Landscape of Spatial Robustness

Spatial Robustness⁷

- Spatial Robustness: rotation, translation
- The neural networks are not robust to spatial attacks
- Attack formation: move a pixel at position (u, v) to

$$\begin{bmatrix} u' \\ v' \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \cdot \begin{bmatrix} u \\ v \end{bmatrix} + \begin{bmatrix} \Delta u \\ \Delta v \end{bmatrix}$$



⁷Engstrom, Logan, et al. "Exploring the Landscape of Spatial Robustness." International Conference on Machine Learning. 2019.

Attack Methods and Defenses

Attacks:

- First Order Methods: the arguments are differentiable
- Grid Search: low dimensional
- Worst of k randomly samples


Defenses:

- Standard Training
- ℓ_∞ -bounded adversarial training
- No random cropping
- Random rotations and translations (30°, 10%pixels)
- Random rotations and translations for larger intervals (40°, 13%pixels)

Different attacks

	MNIST		CIFAR-10		ImageNet	
	Standard	Aug.	Standard	Aug.	Standard	Aug.
Natural	99.31%	99.53%	92.62%	90.02%	75.96%	65.96%
Worst-of-10	73.32%	98.33%	20.13%	79.92%	47.83%	50.62%
First-Order	79.84%	98.78%	62.69%	85.92%	63.12%	66.05%
Grid	26.02%	95.79%	2.80%	58.92%	31.42%	32.90%

Figure 11: Standard vs same level of augmentations



Results

	Model	Nat.	Rand.	Grid	Rand. T.	Grid T.	Rand. R.	Grid R.
MNIST	Standard	99.31%	94.23%	26.02%	98.61%	89.80%	95.68%	70.98%
	ℓ_∞ -Adv	98.65%	88.02%	1.20%	93.72%	34.13%	95.27%	72.03%
	Aug. 30	99.53%	99.35%	95.79%	99.47%	98.66%	99.34%	98.23%
	Aug. 40	99.34%	99.31%	96.95%	99.39%	98.65%	99.40%	98.49%
	W-10 (30)	99.48%	99.37%	97.32%	99.50%	99.01%	99.39%	98.62%
	W-10 (40)	99.42%	99.39%	97.88%	99.45%	98.89%	99.36%	98.85%
CIFAR10	Standard	92.62%	60.93%	2.80%	88.54%	66.17%	75.36%	24.71%
	No Crop	90.34%	54.64%	1.86%	81.95%	46.07%	69.23%	18.34%
	ℓ_∞ -Adv	80.21%	58.33%	6.02%	78.15%	59.02%	62.85%	20.98%
	Aug. 30	90.02%	90.92%	58.90%	91.76%	79.01%	91.14%	76.33%
	Aug. 40	88.83%	91.18%	61.69%	91.53%	77.42%	91.10%	76.80%
	W-10 (30)	91.34%	92.35%	69.17%	92.43%	83.01%	92.33%	81.82%
	W-10 (40)	91.00%	92.11%	71.15%	92.28%	82.15%	92.53%	82.25%
ImageNet	Standard	75.96%	63.39%	31.42%	73.24%	60.42%	67.90%	44.98%
	No Crop	70.81%	59.09%	16.52%	66.75%	45.17%	62.78%	34.17%
	Aug. 30	65.96%	68.60%	32.90%	70.27%	45.72%	69.28%	47.25%
	Aug. 40	66.19%	67.58%	33.86%	69.50%	44.60%	68.88%	48.72%
	W-10 (30)	76.14%	73.19%	52.76%	74.42%	61.18%	73.74%	61.06%
	W-10 (40)	74.64%	71.36%	50.23%	72.86%	59.34%	71.95%	59.23%

Results




Figure 12: Fine-grained analysis

- How to empirically defense? Different ways to obfuscate gradients?
- Certified Robustness: requires good training methods, take long to certify.
- What robustness is reliable and how to achieve is not clear.
- Need more understandings and theories
- Computer Vision vs Natural Language Processing